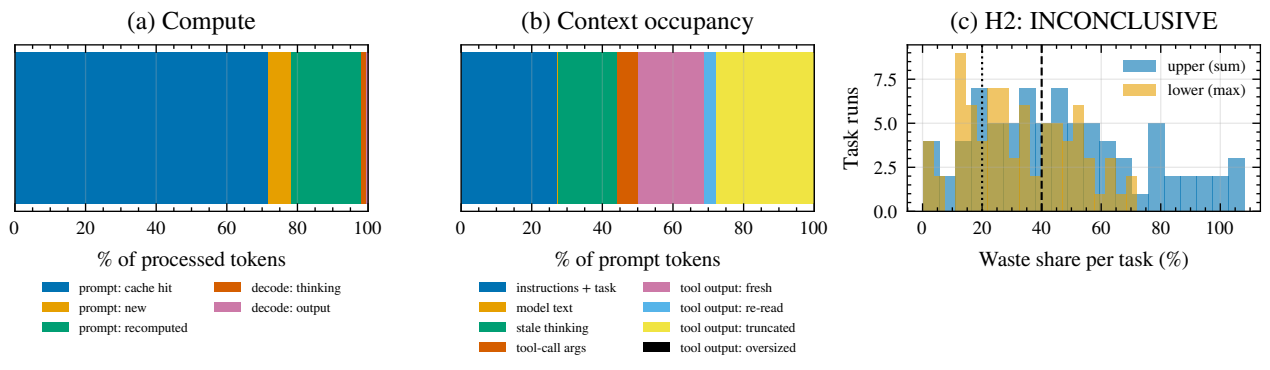
(a) Compute
0
20
40
60
80
100
% of processed tokens
prompt: cache hit
prompt: new
prompt: recomputed
decode: thinking
decode: output
(b) Context occupancy
0
20
40
60
80
100
% of prompt tokens
instructions + task
model text
stale thinking
tool-call args
tool output: fresh
tool output: re-read
tool output: truncated
tool output: oversized
(c) H2: INCONCLUSIVE
Task runs
0.0
2.5
5.0
7.5
upper (sum)
lower (max)
0
20
40
60
80
100
Waste share per task (%)